

I. Description of 100-dimensional embedding.

1. Data preparation

- Downloaded Brown corpus from NLTK and called `brown.words()` return the text in one long list, then made each word lower case.
- Removed stopwords(using `nltk.corpus.stopwords('english')`), punctuations (using `string.punctuation`) and filtered the words which is not alphanumeric, such as " - ", "i'm" and "he's". (only remaining `word.isalnum()`).
- Get a list of 515882 words from corpus after cleaning.

2. Commonly-occurring words

- Used `collections.Counter` function counting the frequencies of words in words list, and applied `most_common()` function to get the words with higher count.
- Select the top 5000 frequent words as vocabulary V , the top 1000 frequent words as context words C , and saved the words and counts as word-count dictionary.

`V_counts_dict: {w: count ...}`

`C_counts_dict: {c: count ...}`

3. Probabilities

- Computed overall distribution $Pr(c)$ of context words, and saved as `pr_c_dict` of length of 1000 of context-word-probability dictionary.

$$Pr(c) = \frac{n(c)}{n(total_c)} ; \# \text{ pr_c_dict: \{c: probability ... \}}$$

- Defined a `window` function to get a surrounding window of four words' indices.

$$Window\ index = [index-2, index-1, index+1, index+2]$$

- Created a nested dictionary called `CW_counts_dict` with value of $n(w, c)$, which can be converted to 5000×1000 matrix.

CW_counts_dict: {w : {c : count}}

Steps:

For each word in cleaned words if the word in vocabulary list(V), check the surrounding window of four words, then for each surrounding words, if the word in context words list(C), add 1 to the corresponding CW_counts_dict by key as [w][c].

- Divided CW_counts_dict[w][c] values by V_counts_dict[w] value to get $Pr(c|w)$, saved as a nested dictionary called **pr_cw_dict**.

$Pr(c|w) = \frac{n(w,c)}{n(w)}$; # pr_cw_dict: {w : {c : probability}}

4. $\phi(w)$

Represented each vocabulary item by a $|C|$ -dimensional vector $\phi(w)$, that each row represents each vocabulary w by 1000 dimensional vector. Define a function to get $\phi(w)$. $\phi(w) = \max(0, \log \frac{Pr(c|w)}{Pr(c)})$

Steps:

- Calculated $\frac{Pr(c|w)}{Pr(c)}$, $\phi(w)[c] = \max(0, \log \frac{Pr(c|w)}{Pr(c)})$
- If $\frac{Pr(c|w)}{Pr(c)} > 1$, chose $\log \frac{Pr(c|w)}{Pr(c)}$, else chose 0, then saved as another nested dictionary(**phi**). # phi: {w: {c: $\phi(w)$... } ... }
- Converted the nested dictionary of $\phi(w)$ to a dataframe.
phi_df = pd.DataFrame(phi).T

5. 100 dimensional representation

Using sklearn.decomposition.PCA to do PCA on $\phi(w)$ transformed 1000-dimension to 100-dimension, saved as a dataframe called **phi_transformed_df**. $\Psi(w) \in R^{100}$

```
pca = PCA(n_components=100)
phi_transformed = pca.fit_transform(phi_df)
phi_transformed_d = pd.DataFrame(phi_transformed, index=V)
```

II. Nearest Neighbor results.

1. Cosine_distance:

Defined a function of find nearest neighbor. Using `sklearn.metrics.pairwise_distances()` calculated the distance between given vocabulary(w) and each word ($w' \neq w$) in cleaned text, and saved w' -distance pair dictionary, then from the dictionary get the key of min value. `NearestNeighbor = min(dict, key=dict.get)`

- Cosine distance: (set `metric='cosine'`)

Nearest neighbor: communism – utopian

Nearest neighbor: autumn – slide

Nearest neighbor: cigarette – smelled

Nearest neighbor: pulmonary – artery

Nearest neighbor: mankind – civil

Nearest neighbor: africa – asia

Nearest neighbor: chicago – portland

Nearest neighbor: revolution – exercise

Nearest neighbor: september – december

Nearest neighbor: chemical – milligrams

Nearest neighbor: detergent – cleaning

Nearest neighbor: dictionary – text

Nearest neighbor: storm – autumn

Nearest neighbor: worship – religion

Nearest neighbor: school – college

Nearest neighbor: weather – hot

Nearest neighbor: tax – income

Nearest neighbor: art – literature

Nearest neighbor: data – results

Nearest neighbor: america – states

Nearest neighbor: earth – dust

Nearest neighbor: happy – know

Nearest neighbor: female – feature

Nearest neighbor: color – dress

Nearest neighbor: flight – landing

- Euclidean distance: (set `matric = 'euclidean'`)

Nearest neighbor: communism – utopian

Nearest neighbor: autumn – slide

Nearest neighbor: cigarette – smelled

Nearest neighbor: pulmonary – artery

Nearest neighbor: mankind – fatal

Nearest neighbor: africa – asia

Nearest neighbor: chicago – portland

Nearest neighbor: revolution – twentieth

Nearest neighbor: september – december

Nearest neighbor: chemical – milligrams

Nearest neighbor: detergent – fabrics

Nearest neighbor: dictionary – text

Nearest neighbor: storm – autumn

Nearest neighbor: worship – voluntary

Nearest neighbor: school – college

Nearest neighbor: weather – cheap

Nearest neighbor: tax – income

Nearest neighbor: art – english

Nearest neighbor: data – include

Nearest neighbor: america – peoples

Nearest neighbor: earth – dust

Nearest neighbor: happy – wondered

Nearest neighbor: female – nude

Nearest neighbor: color – pure

Nearest neighbor: flight – landing

As we can see from the results, the nearest neighbor words are have high similarity, such as 'school' – 'college', 'september' – 'december', 'tax' – 'income', and 'chicago – portland', and there's no much difference between the two distance measurments.

III. Clustering

Tried to use two different algorithm, **KMeans** and **Agglomerative** clustering, and using **Silhouette Coefficient** to compared different model that if the clusters dense and well seperated.

- **KMeans**: `sklearn.cluster.KMeans` Separate samples in n groups of equal variance, minimizing within-cluster sum-of-squares, using euclidean distance.
`k_means = KMeans(n_clusters=100,init='k-means++')`
`silhouette_score= -0.017`
- **Hierarchical clustering**: `sklearn.cluster.AgglomerativeClustering` Set affinity as 'cosine distance' and 'euclidean distance' with four different linkage. It supports Ward, single, average, and complete linkage strategies.
 - **ward** minimizes the variance of the clusters being merged. Only for euclidean distance.
 - **average** uses the average of the distances of each observation of the two sets.
 - **complete** or maximum linkage uses the maximum distances between all observations of the two sets.
 - **single** uses the minimum of the distances between all observations of the two sets.

Fitted different clustering model on 100-dimensional transformed $\phi(w)$ data, and checked each model silhouette coefficient, then checked each cluster word count with high silhouette_score. After compareration, I found that even though some of the Agglomerative clustering model like (`affinity='euclidean',linkage='complete'`) have high silhouette score, most of the clusers have little words, which caused by overfitting.

Since the `AgglomerativeClustering(affinity='euclidean',linkage='ward')` is same as the `KMeans` algorithm, and `KMeans` silhouette_score is higher than `AgglomerativeClustering` model, I selected `KMeans` algorithm and euclidean distance for the clustering. And the clusters seems coherent.

A few of the meaningful clusters:

- **group_12** : ['school', 'children', 'members', 'college', 'university', 'students', 'class', 'schools', 'student', 'activities', 'negro', 'interested', 'professional', 'parents', 'active', 'jewish', 'teachers', 'teaching']

- **group_55** : ['cost', 'total', 'rate', 'tax', 'increase', 'costs', 'pay', 'amount', 'higher', 'increased', 'due', '1961', 'sales', 'average', 'lower', 'additional', 'fiscal', 'income', 'rates', 'annual', 'operating', 'gain', 'estimated', 'wages']
- **group_64**: ['white', 'black', 'red', 'dark', 'brown', 'hair', 'blue', 'color', 'beautiful', 'arms', 'blood', 'green', 'sun', 'heavy', 'deep', 'mouth', 'teeth', 'thin', 'bright', 'rose', 'neck', 'watched', 'gray', 'walls', 'dress', 'thick', 'wore', 'nose', 'stared', 'sky', 'pale', 'tall', 'yellow', 'tiny', 'pink', 'skin', 'golden', 'shade', 'beard', 'blonde']
- **group_80**: ['war', 'south', 'west', 'north', 'peace', 'america', 'east', 'europe', 'british', 'germany', 'france', 'berlin', 'russia', 'china', 'britain', 'africa', 'asia', 'atlantic', 'eastern']
- **group_93** : ['day', 'last', 'year', 'later', 'next', 'days', 'early', 'week', 'morning', 'months', 'return', 'late', 'hours', 'weeks', 'nearly', 'summer', 'month', 'spring', 'spent']
- **group_99** : ['years', 'three', 'four', 'five', 'ago', 'six', 'minutes', 'miles', 'hundred', 'ten', 'seven', 'eight', 'dollars', 'thousand', 'nine', 'twenty', 'fifty', 'thirty', 'fifteen', 'twelve', 'eleven']